

Decision Trees and Data Structures

Communication Complexity Seminar


Instructor: Ronitt Rubinfeld

Fall 2009

By: Ayelet Leigh


Plan for today:

- communication complexity lower bounds yield data structure lower bounds



What are we doing ?

- Implementation of a “database” - D :
 - D represents a subset S of $\{1 \dots N\}$






- Access to D via "membership queries" - Q
 - for each i , can ask: “ $i \in S$?”
 - denote possible queries by $Q = \{1, \dots, N\}$

The Cell Probe Model

check,
investigate,
inquiry

- Database Definition:
 - Memory mapped into fixed-size cells (b bits per cell)
- Query via Decision Trees
 - Each node probes one cell
 - Continue based on bits received (2^b options)
 - Leaves contain query answers
- Time = $\max_{\#probs}$ Depth of worst decision tree

Read value of



What is the goal today?

- Won't see *time* lower-bound
- Won't see *space* lower-bound
- Will see *trade-off* lower-bound
 - Time-space trade-off
- Reduction from database problem to communication complexity problem
 - We know communication lower-bounds



Reduction to communication complexity problem

- Function $f(q, D)$ – answer to q on D :
 - Assumptions:
 - D is stored using s cells of b bits each
 - Any $q \in Q$ can be solved using at most t probes
 - Lemma 1: f can be solved by t -round communication protocol, at each round:
 - Alice sends $\log(s)$ bits holds q
 - Bob sends b bits holds D

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes

To Show
Each of $\leq t$ rounds:
Alice $\log s$ bits
Bob b bits

Proof

Alice

Bob

Simulates the tree corresponding q
Implements probe to cell by sending cell index

There are s cells



$\log(s)$ bits

Reply with cell content



b bits

Updates her current node accordingly



There are only t probes

Example 1

- U – vector subspace of \mathbb{Z}_2^n
- $\forall q \in \{0,1\}^n$ We need to answer: " $q \in U$?"
- Cell size: $b = n$ bits
- How to store subspace U ?
 - Set of $n - \dim(U)$ linear equations defining U

Reminder...

U subspace of Z_2^n

$$n = 2$$

Linear equations defining subspace U :

$$(x \ y) \begin{pmatrix} 1 \\ 1 \end{pmatrix} = 0 \implies U: (1 \ 1), (0 \ 0)$$

$$\dim(U) = 1 \implies 1 \text{ equation}$$

$$\dim(U) = 2 \implies 0 \text{ equations} - \text{all } z_2^2$$

$$(x \ y) \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} = 0 \implies U: (0 \ 0)$$

$$\dim(U) = 0 \implies 2 \text{ equation}$$


Example 1

- U – vector subspace of Z_2^n
- $\forall q \in \{0,1\}^n$ We need answer: " $q \in U$?"
- Cell size: $b = n$ bits
- How to store subspace U ?
 - Set of $n - \dim(U)$ linear equations defining U
 - Use cell for each equation
 - Reminder: every equation in U described by n bits
 - $\forall q$ we access all cells to check if q satisfies each of them

Space: $s=n$ cells

Query Time: $t=n$


Example - can we do it faster?

- Can we answer with $t = o(n)$ time?
 - How many cells needed?
 - Reminder: $f(q, U)$ is the SPAN problem (last week)
 - We saw: Alice sends $\Omega(n)$ or Bob sends $\Omega(n^2)$
 - From lemma 1, using s -cell solution for subspace gives protocol for SPAN, satisfying:
 - Alice sends $t \log s$ bits
 - Bob sends tn bits
 - If $t = o(n)$:
 - Then Bob sends $o(n^2)$ bits
 - Therefore Alice sends $\Omega(n)$ bits
-  $s \geq 2^{\Omega(n/t)}$



Dynamic Data Structures


- Definition:
 - A data structure whose data may change during its lifetime.
- The problem:
 - Maintaining a database D under update operations:



The dynamic problem

- For any k , the dynamic problem function:
$$f_k(q, \vec{u})$$
 - $\vec{u} = u_1, \dots, u_k$ - sequence of updates
 - q - a query
- The value of f is the answer to the query, after the updates were performed.

Time lower bound



- Assumptions:
 - Every **update** can be performed in $\leq t$ probes
 - Every **query** can be performed in $\leq t$ probes
 - Cell size is **b bits**
- Lemma 2:
 f_k has a communication protocol where:
 - **Alice** sends $O(t \log(kt))$ bits
 - **Bob** sends $O(t(b + \log|Q|) + \log(kt))$ bits

Notice similarity to lemma 1

The proof will be similar too

Before we prove

- More assumptions:
 - Both Alice and Bob know the initial DB
 - Only Bob knows the updates
- Reminder - hash functions:
 - mod function reduces the number of bits used

Useful fact

W is a subset of $\{1, \dots, s\}$, $|W| \leq kt$

Random prime $p \in 1, \dots, |W|^3$ will satisfy

$$\forall w, w' \in W, w \not\equiv w' \pmod{p}$$

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes
Updates made with $\leq t$ probes

To Show
Each of $\leq t$ rounds:
Alice $O(t \log(kt))$ bits
Bob $O(t(b + \log|Q|) + \log(kt))$ bits

Proof

Alice

Bob

Simulates the updates: u_1, \dots, u_k

Fix W - set of cells that changed during updates

$$|W| \leq kt$$

Find a prime $p \leq |W|^3$ s.t.

$$\forall w, w' \in W, w \not\equiv w' \pmod{p}$$

Sends p



$\log(kt)$ bits

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes
Updates made with $\leq t$ probes

To Show
Each of $\leq t$ rounds:
Alice $O(t \log(kt))$ bits
Bob $O(t(b + \log|Q|) + \log(kt))$ bits

Proof – cont.

Alice



Bob

Simulates the tree corresponding q
Implements probe to cell m by:

Sending $m \bmod p$



$O(\log(kt))$
bits

Finds unique $w \in W, w = m \bmod p$

Sends
address and the content of w



Address: $\log s$ bits
Content: b bits

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes
Updates made with $\leq t$ probes

To Show
Each of $\leq t$ rounds:
Alice $O(t \log(kt))$ bits
Bob $O(t(b + \log|Q|) + \log(kt))$ bits

Proof – cont.

Alice

Bob

If $w = m$:

Alice has the content of m that
Bob sent

If $w \neq m$:

The updates didn't change cell m
Alice has the initial value of m



There are only t probs

$O(t * \log(kt))$
bits



$O(t(b + \log s) + \log(kt))$
bits

prime p
(sent once)

every loop

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes
Updates made with $\leq t$ probes

To Show
Each of $\leq t$ rounds:

Bob $O(t(b + \log|Q|) + \log(kt))$ bits

Proof - finish

- This complexity is small if:
 - s is known and small enough.
- But not necessarily if $s > |Q|$
 - Can happen: DB must keep answers to all queries, and be able to answer after future updates.
- The use of s not useful to get time lower bound!

Assumptions:
 s cells b bits per cell
 q solves with $\leq t$ probes
Updates made with $\leq t$ probes

Trick for saving bits

Bob

Uses encoding for w with $\log|Q|$ bits:

Conclude what q' is from the previous probes

If all correct the q' Bob found matches Alices

Sends $b + \log|Q|$ bits



$O(t(b + \log|Q|) + \log(kt))$
bits

Alice

If $w = m$:

$q = q'$

Alice has the content of m that
Bob sent

If $w \neq m$:

The updates didn't change cell m

Alice has the initial value of m



Example 2

- U – vector subspace of Z_2^n
- U can be updated using $add(u)$ operation
 - $add(u)$: replaces U by $span(U \cup \{u\})$
- Single query: $\dim(U)$
- Cell size: $b = n$ bits
- $\forall q \in \{0,1\}^n$ We need to answer: " $q \in U$?"

Example 2 – lower bound

- Lower bound for the problem " $q \in U?$ "
 - Reminder: $f(q, U)$ is the SPAN problem (last week)
 - We saw: Alice sends $\Omega(n)$ or Bob sends $\Omega(n^2)$
 - From lemma 2, using s-cell solution for subspace gives protocol for SPAN, satisfying:
 - Alice sends $o(t \log kt)$ bits
 - Bob sends $o(t(b + \log|Q|) + \log kt)$ bits
 - In both cases, by using lemma 2, When $|Q| = 2^n$



$$t \log kt = \Omega(n)$$

$$t(b + \log|Q|) + \log kt = \Omega(n^2)$$

number of probes needed :

$$t = \Omega(n / \log(n))$$

$$U \subseteq Z_2^n$$

$$\text{add}(u) = \text{span}(U \cup \{u\})$$

$$\text{Query} = \dim(U)$$

$$b = n \text{ bits}$$



Questions
are
guaranteed in
life;
Answers
aren't.



That's all Folks!